


GPS-conditioned JEPA pretraining and downstream reuse

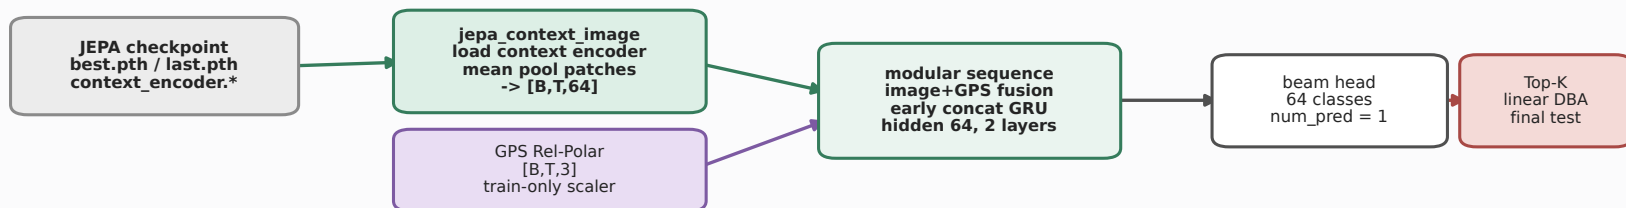
Self-supervised image token prediction guided by GPS relative-polar geometry; only the context image encoder is reused for supervised beam prediction.

a JEPA pretraining path



Pretraining excludes beam labels, legacy KD distillers, and frozen external teacher checkpoints.

b Downstream reuse for supervised image+GPS beam prediction



Current 2604/fair configs reuse S32-S34 JEPA checkpoints; scene31-only checkpoints are excluded from fair defaults.